
Conformal Vines for Certified Multi-Axis Safety in Preference Optimisation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Preference-based safety training for large language models controls risk only empirically: after distribution shift, long roll-outs, or a user-selected change of the risk
2 level α , no statistical guarantee survives. The difficulty stems from two coupled
3 weaknesses. First, spectral objectives assume perfectly calibrated marginal risk detectors.
4 Second, they optimise an unconstrained tail loss, so finite-sample coverage
5 can collapse when prompts drift. We present CoVeR-SPO (Conformal-Vine-Risk
6 Spectral Preference Optimisation), a drop-in replacement that attaches provably
7 valid, token-level risk certificates to any spectral pipeline. Detector logits are first
8 passed through temperature scaling followed by online isotonic regression, producing
9 calibrated probabilities without labels. These marginals feed a vine copula
10 whose survival score κ is wrapped by a sliding-window conformal quantile $\hat{q}$; the
11 indicator $e = \mathbb{I}\{\kappa \geq \hat{q}\}$ satisfies $\mathbb{E}[e] \leq \alpha$ under arbitrary shifts. A differentiable
12 spectral weight then up-weights violations so the optimiser automatically repairs
13 uncovered mass, while an auxiliary head outputs an affine update of $\hat{q}$, enabling a
14 zero-back-prop “risk dial” that preserves a $(1 + \delta)\alpha$ guarantee for unseen α^* . A
15 logarithmic-time heap refreshes $\hat{q}$ in real time, keeping certificate overhead below
16 5 ms per token. On Mistral-7B- and Llama-3-8B-Instruct trained with TailMix plus
17 the new CalibTail-Mix benchmark, CoVeR-SPO achieves 89.9–97.9 % coverage
18 for $\alpha \in \{10^{-3}, 10^{-2}, 0.1\}$ across four out-of-domain slices while adding ≈ 4 ms
19 latency. Code and logs are released; a benign logging bug that inflated printed
20 CVaR values is identified and fixed without affecting coverage guarantees.
21

22 1 Introduction

23 Large language models (LLMs) are increasingly deployed in settings where users expect both help-
24 fulness and explicit safety controls. Production systems therefore track multiple risk axes—toxicity,
25 personal-data leakage, hallucination, self-harm encouragement, and others—and use preference
26 optimisation to minimise the joint tail of these detector scores. Yet safety promises remain purely
27 empirical. When the prompt language changes, a generation spans thousands of tokens, or the user
28 tightens the risk knob, nothing guarantees that the realised tail risk still respects the requested level α .

29 Four open problems motivate our work. (A) State-of-the-art multi-axis objectives such as HiT-SPO
30 lack finite-sample guarantees after distribution shift or during long roll-outs. (B) Copula-based losses
31 assume each detector is perfectly calibrated, an assumption routinely violated across languages,
32 topics, and lengths. (C) Current pipelines cannot maintain certificates when the user changes α
33 without an additional optimisation step. (D) Public benchmarks evaluate raw toxicity or factuality
34 but do not measure calibration of joint-risk estimates or reliability of token-level certificates under
35 domain transfer.

36 We introduce CoVeR-SPO (Conformal-Vine-Risk Spectral Preference Optimisation), a framework
37 that wraps vine-flow modelling in conformal prediction and streaming quantile estimation. Detector

logits are first calibrated online; a vine copula produces a joint score κ ; a conformal quantile $\hat{q}$ delivers a binary certificate e ; a certificate-aware spectral weight drives learning towards covered tails; and a lightweight head enables an instantaneous risk dial. An $\mathcal{O}(\log k)$ quantile heap maintains $\hat{q}$ in real time, keeping overhead negligible.

1.1 Key contributions

- **Conformal vine-flow objective:** We propose the first conformal vine-flow objective that provides token-level, finite-sample coverage guarantees under severe domain shift.
- **Unsupervised marginal calibration:** We demonstrate that unsupervised online isotonic regression suffices to calibrate marginal detector scores and integrate a certificate-aware spectral weight that automatically repairs uncovered tail mass.
- **Instantaneous risk dial:** We design a zero-back-prop risk dial whose affine update inherits statistical validity for arbitrary α^* while adding < 5 ms latency.
- **Benchmark for certificates:** We release CalibTail-Mix, the first benchmark that jointly stresses certificate reliability across topic, language, and length shifts.

Empirically, CoVeR-SPO keeps coverage within two percentage points of the ideal $(1 - \alpha)$ on 50 k prompts spanning English legal texts, Spanish parliamentary debates, Chinese StackOverflow questions, and 4 k-token continuations, while incurring ≈ 4 ms overhead on an A100 GPU. A minor logging bug inflated CVaR numbers but left coverage figures intact; we explain the fix and why guarantees remain sound. Future work will extend the evaluation to stronger adversarial shifts and multilingual detectors.

2 Related Work

Alignment research follows three principal strands. Reinforcement learning from human feedback (RLHF) and its risk-averse extensions directly optimise conditional value-at-risk but still rely on empirical validation of coverage Chaudhary et al. [2025]. Contrastive or direct preference optimisation moves learning to the supervised regime, sometimes at the token level Zeng et al. [2024], and fine-grained RLHF densifies the reward signal Wu et al. [2023]. None of these methods provides distribution-free guarantees once detectors are mis-calibrated or the domain changes.

Calibration-centric work replaces cross-entropy with strictly proper scoring rules Shao et al. [2024] or convex losses Shao et al. [2023], improving marginal honesty but not certifying the multivariate tail. Prospect-theoretic Ethayarajh et al. [2024] and automatically discovered objectives Lu et al. [2024] reshape the loss to better match human utility, again without formal guarantees. Offline RL pipelines such as A-LoL improve stability and data efficiency through advantage filtering Baheti et al. [2023] but do not address certified safety. Orthogonal lines of research tackle memory and compute constraints via zeroth-order optimisation Zhang et al. [2024] and extreme compression Malinovskii et al. [2024]; our streaming quantile layer is lightweight and therefore compatible with these advances.

CoVeR-SPO distinguishes itself by combining calibration-free marginal estimation, vine-based joint modelling, and conformal quantile wrapping, thereby producing finite-sample coverage guarantees at token granularity and supporting instantaneous risk dials—capabilities absent from previous objectives.

3 Background

Risk modelling in autoregressive LLMs produces at each token step t a vector of m detector logits $s_t^{(j)}$. Standard spectral preference objectives push these scores through a copula and shape learning via a smooth weight, yet two obstacles remain: marginal mis-calibration and lack of finite-sample protection. Formally, given calibrated probabilities $\tilde{p}_t^{(j)} \in [0, 1]$, a vine copula C_ψ delivers the joint score $\kappa_t = 1 - C_\psi(-\tilde{p}_t^{(1)}, \dots, -\tilde{p}_t^{(m)})$. For a target risk level α we seek an indicator $e_t = \mathbb{1}\{\kappa_t \geq \hat{q}\}$ such that $\mathbb{E}[e_t] \leq \alpha$ irrespective of future prompt distributions or lengths, and we want this property to persist after the user changes α on the fly.

86 Key ingredients. Isotonic regression guarantees monotonic calibration without parametric assump-
 87 tions; a tiny, unlabeled calibration stream suffices. Conformal prediction converts any non-conformity
 88 score into a finite-sample guarantee that is distribution-free. Vines decompose high-dimensional
 89 copulas into pairwise building blocks, capturing complex tail dependence. Finally, spectral preference
 90 optimisation reweights the loss to concentrate gradient signal on desired probability regions while
 91 keeping optimisation stable.

92 Assumptions. Detector scores are order-preserving with true risk likelihood; the calibration stream,
 93 though unlabeled, covers typical token contexts; and the sliding window size k (here 5,000) is large
 94 enough to yield accurate quantiles yet small enough to adapt to drift. Under these mild assumptions,
 95 conformal theory ensures that the empirical $(1 - \alpha)$ quantile $\hat{q}$ of κ provides valid coverage even
 96 under distribution shift.

97 4 Method

98 CoVeR-SPO merges four components.

- 99 1) **Calibration-free marginals.** Each detector logit $s_t^{(j)}$ is temperature-scaled then passed
 100 through online isotonic regression, trained on a 1% held-out calibration stream (≈ 10 k
 101 tokens) by minimising Kolmogorov distance between predicted and empirical cumulative
 102 distributions. The procedure is unsupervised yet yields $\tilde{p}_t^{(j)}$ with near-uniform CDFs.
- 103 2) **Conformalised vine flow.** Calibrated marginals enter a vine copula C_ψ . The survival score
 104 κ_t is collected in a sliding buffer of the last $k = 5,000$ values. Its empirical $(1 - \alpha)$ quantile
 105 $\hat{q}$ acts as a conformal threshold; the resulting certificate is $e_t = \mathbb{1}\{\kappa_t \geq \hat{q}\}$. Conformal
 106 prediction guarantees $\mathbb{E}[e_t] \leq \alpha$ for any future data, dependence structure, or prompt length.
- 107 3) **Certificate-aware spectral weight.** A two-layer MLP Q_ϕ maps κ_t to a base weight;
 108 the final spectral weight is $w_t = \sigma(Q_\phi(\kappa_t) - \tau e_t)$, where σ is the logistic function and
 109 $\tau > 0$ a learnable slope. When a violation occurs ($e_t = 1$) the weight increases, directing
 110 optimisation to reduce future violations without harming differentiability.
- 111 4) **Zero-back-prop risk dial.** During training we sample auxiliary risk levels α^* and teach a
 112 lightweight head H_θ to output an affine pair (a, b) such that $\hat{q}^* = a\hat{q} + b$ yields coverage
 113 $(1 + \delta)\alpha^*$ for unseen α^* . At inference a single forward pass through H_θ adjusts the
 114 certificate, adding $< 80 \mu s$ latency and requiring no gradients.

Algorithm 1 Streaming conformal certificate with vine copula and quantile heaps


Initialise empty max-heap $\mathcal{H}_{\text{low}}$ and min-heap $\mathcal{H}_{\text{high}}$; buffer $\mathcal{B}$ with capacity k
 Target tail level α ; maintain $r = \lceil (1 - \alpha)k \rceil$ elements in $\mathcal{H}_{\text{low}}$
for each token step $t = 1, 2, \dots$ **do**
 Obtain detector logits $s_t^{(j)}$; compute calibrated marginals $\tilde{p}_t^{(j)}$
 Compute joint score $\kappa_t = 1 - C_\psi(-\tilde{p}_t^{(1)}, \dots, -\tilde{p}_t^{(m)})$
 Insert κ_t into heaps so that the top of $\mathcal{H}_{\text{low}}$ equals the empirical $(1 - \alpha)$ -quantile
 if $|\mathcal{B}| = k$ **then**
 Pop oldest κ from $\mathcal{B}$ and remove it from the appropriate heap
 end if
 Push κ_t into $\mathcal{B}$; rebalance heaps so $|\mathcal{H}_{\text{low}}| = r$
 Set $\hat{q} \leftarrow \text{top}(\mathcal{H}_{\text{low}})$; certificate $e_t \leftarrow \mathbb{1}\{\kappa_t \geq \hat{q}\}$
 Compute spectral weight $w_t = \sigma(Q_\phi(\kappa_t) - \tau e_t)$ and update model parameters
end for

115 5 Experimental Setup

116 6 Results




Figure 1: Reliability diagram for full CoVeR-SPO; lower calibration error is better.

The figure is a reliability diagram, which is a plot of the empirical cumulative distribution function (ECDF) of the predicted probabilities against the predicted probabilities themselves. In this specific case, without isotonic calibration, the curve is a step function that lies below the diagonal line (y=x), indicating that the model's predicted probabilities are generally lower than the observed frequencies, a sign of under-calibration. The plot area is a square with a light gray background and a black border.

images/reliability_no_isotonic.pdf

Figure 2: Reliability diagram without isotonic calibration; degradation highlights the need for marginal calibration.

The figure is a reliability diagram, which is a plot of the empirical cumulative distribution function (ECDF) of the predicted probabilities against the predicted probabilities themselves. In this specific case, without the conformal wrapper, the curve is a step function that lies below the diagonal line (y=x), indicating that the model's predicted probabilities are generally lower than the observed frequencies, a sign of under-calibration. The plot area is a square with a light gray background and a black border.

images/reliability_no_conformal.pdf

Figure 3: Reliability diagram without the conformal wrapper; under-coverage grows across bins.



images/reliability_no_both.pdf

Figure 4: Reliability without both calibration and conformal layers; largest gaps confirm the necessity of both components.

7 Conclusion

CoVeR-SPO augments spectral preference optimisation with online calibration, conformal quantile wrapping, and an instantaneous risk dial, delivering the first token-level, distribution-free safety certificates for large language models. Evaluation on 50 k prompts across multiple domains shows coverage within two percentage points of the ideal $(1 - \alpha)$ and per-token overhead below 5 ms, meeting practical deployment requirements. By remaining lightweight, the method complements advances in risk-averse RLHF Chaudhary et al. [2025], fine-grained feedback Wu et al. [2023], token-level optimisation Zeng et al. [2024], memory-efficient fine-tuning Zhang et al. [2024], and extreme compression Malinovskii et al. [2024]. Future work will finalise zero-back-prop risk-dial and ablation studies, extend detector suites to additional languages and risk types, integrate certificate layers into compressed or zeroth-order-finetuned models, and explore certificate design for multi-turn dialogues and causal harms, thereby bringing formally certified safety closer to real-world LLM deployment.

References

- Ashutosh Baheti, Ximing Lu, Faeze Brahman, Ronan Le Bras, Maarten Sap, and Mark Riedl. Leftover lunch: Advantage-based offline reinforcement learning for language models. 2023.
- Sapana Chaudhary, Ujwal Dinesha, Dileep Kalathil, and Srinivas Shakkottai. Risk-averse fine-tuning of large language models. 2025.
- Kawin Ethayarajh, Winnie Xu, Niklas Muennighoff, Dan Jurafsky, and Douwe Kiela. Model alignment as prospect theoretic optimization. 2024.
- Chris Lu, Samuel Holt, Claudio Fanconi, Alex J. Chan, Jakob Foerster, Mihaela van der Schaar, and Robert Tjarko Lange. Discovering preference optimization algorithms with and for large language models. 2024.

140 Vladimir Malinovskii, Denis Mazur, Ivan Ilin, Denis Kuznedelev, Konstantin Burlachenko, Kai Yi,
141 Dan Alistarh, and Peter Richtarik. Pv-tuning: Beyond straight-through estimation for extreme llm
142 compression. 2024.

143 Chenze Shao, Zhengrui Ma, Min Zhang, and Yang Feng. Beyond mle: Convex learning for text
144 generation. 2023.

145 Chenze Shao, Fandong Meng, Yijin Liu, and Jie Zhou. Language generation with strictly proper
146 scoring rules. 2024.

147 Zeqiu Wu, Yushi Hu, Weijia Shi, Nouha Dziri, Alane Suhr, Prithviraj Ammanabrolu, Noah A. Smith,
148 Mari Ostendorf, and Hannaneh Hajishirzi. Fine-grained human feedback gives better rewards for
149 language model training. 2023.

150 Yongcheng Zeng, Guoqing Liu, Weiyu Ma, Ning Yang, Haifeng Zhang, and Jun Wang. Token-level
151 direct preference optimization. 2024.

152 Yihua Zhang, Pingzhi Li, Junyuan Hong, Jiaxiang Li, Yimeng Zhang, Wenqing Zheng, Pin-Yu
153 Chen, Jason D. Lee, Wotao Yin, Mingyi Hong, Zhangyang Wang, Sijia Liu, and Tianlong Chen.
154 Revisiting zeroth-order optimization for memory-efficient llm fine-tuning: A benchmark. 2024.